



COMP2017 COMP9017

Assignment 2

Due: 11:59PM Tuesday 28 March 2023 local Sydney time

This assignment is worth 5% + 30% of your final assessment

Task Description

Ah, yes, because what the world really needs right now is yet another virtual machine. But not just any virtual machine, no! We need one with the highly coveted and incredibly useful feature of heap banks. Because who needs to worry about memory allocation when you can just throw it in a heap, am I right? So gather 'round, folks, and get ready to develop the `vm_RISKXVII` because nothing says "cutting edge" like a virtual machine named after a board game that this assignment has absolutely nothing to do with. Now, let's dive into the specs and get this party started!

In this assignment you will be implementing a simple virtual machine. Your program will take a single command line argument being the path to the file containing your RISK-XVII assembly code.

Before attempting this assignment it would be a good idea to familiarise yourself with registers, memory space, program counters, assembly and machine code. A strong understanding of these concepts is essential to completing this assignment. Section 3.6 and 3.7 of the course textbook provide specific detail to x86_64 architecture, however you can review these as a reference.

In order to complete this assignment at a technical level you should revise your understanding of bitwise operations, file IO, pointers and arrays.

Some implementation details are purposefully left ambiguous; you have the freedom to decide on the specifics yourself. Additionally this description does not define all possible behaviour that can be exhibited by the system; some error cases are not documented. You are expected to gracefully report and handle these errors yourself.

You are encouraged to ask questions on Ed¹. Make sure your question post is of "**Question**" post type and is under "**Assignment**" category → "**A2**" subcategory. As with any assignment, make sure that your work is your own², and that you do not share your code or solutions with other students.

The Architecture

In this assignment you will implement a virtual machine for an 32-bit instruction-set. The memory mapped virtual components of your machine are outlined below:

¹<https://edstem.org/au/courses/10466/discussion/>

²Not GPT-3/4's, ChatGPT's or copilot's, etc.

- **0x0000 - 0x3fff**: Instruction Memory - Contains 2^{10} of bytes for text segment.
- **0x0400 - 0x7fff**: Data Memory - Contains 2^{10} of bytes for global variables, and function stack.
- **0x0800 - 0x8fff**: Virtual Routines - Accesses to these address will cause special operations to be called.
- **0xb700 +**: Heap Banks - Hardware managed 128 x 64 bytes banks of dynamically allocate-able memory.

Your machine also has a total of 32 registers, as well as a PC (program counter) that points to the address of the current instruction in memory. Each of the general-purpose registers can store 4 bytes (32 bits) of data that can be used directly as operands for instructions. **All registers are general-purpose except for the first one**, which has an address of 0. This register is called the **zero** register, as any read from it will return a value of 0. Writes to the **zero** register are ignored.

During execution you should not store any information about the state of the machine outside of the virtual memory devices and the register bank.

Note: A register stores a single value using a fixed bit width. It the size of a register corresponding to the processor *word size*, in this case 32 bits. Think of them as a primitive variable. Physical processor hardware is constrained, and the number of registers is always fixed. There are registers which serve specific purposes, and those which are general. Please identify these in the description and consider them for your solution. You need not consider special purpose registers, such as floating point, in this assignment.

<https://tutorcs.com>

RISK-XVII Instruction Set Architecture

An Instructions-Set Architecture(ISA) specifies a set of instructions that can be accepted and executed by the target machine. A program for the target machine is an ordered sequence of instructions.

Our virtual machine will operate on a home-brewed 'RISK-XVII' instruction set architecture. During marking, you will be provided with binaries in this ISA to run on your virtual machine. RISK-XVII is a reduced version of the well-known RV32I instruction set architecture, and your virtual machine should be able to execute binary programs compiled for RV32I, as long as they do not include instructions that were not specified by 'RISK-XVII'.

There are in total 33 instructions defined in RISK-XVII, they can be classified into three groups by their functionality:

1. Arithmetic and Logic Operations - e.g. `add`, `sub`, `and`, `or`, `slt`, `sll`
2. Memory Access Operations - e.g. `sw`, `lw`, `lui`
3. Program Flow Operations - e.g. `jal`, `beq`, `blt`

These instructions provide access to memory and perform operations on data stored in registers, as well as branching to different locations within the program to support conditional execution. Some instructions also contain data, i.e., an immediate number, within their encoding. This type of instruction is typically used to introduce hard-coded values such as constants or memory address offsets.

The RISK-XVII instruction set is Turing complete and, therefore, can run any arbitrary program, just like your PC!

Instructions in the RISK-XVII instruction set architecture are encoded into 4 bytes of data. Since each instruction can access different parts of the system, seven types of encoding formats were designed to best utilize the 32 bits of data to represent the operations specified by each instruction: R, I, S, SB, U, UJ. The exact binary format of each encoding type can be found in the table below.:

Type	Format												
	31	25	24	20	19	15	14	12	11	7	6	0	
R	func7			rs2		rs1		func3		rd		opcode	
I	imm[11:0]				rs1		func3		rd		opcode		
S	imm[11:5]			rs2		rs1		func3		imm[4:0]		opcode	
SB	imm[12 10:5]			rs2		rs1		func3		imm[4:1 11]		opcode	
U	imm[32:12]								rd		opcode		
UJ	imm[20 10:1 11 19:12]								rd		opcode		

Let's take a look at some common fields in all types of encoding:

- **opcode** - used in all encoding to differentiate the operation, and even the encoding type itself.
- **rd, rs1, rs2** - register specifiers. **rs1** and **rs2** specify registers to be used as the source operand, while **rd** specifies the target register. Note that since there are 32 registers in total, all register specifiers are 5 bits in length.
- **func3, func7** - these are additional opcodes that specify the operation in more detail. For example, all arithmetic instructions may use the same **opcode**, but the actual operation, e.g. add, logic shift, are defined by the value of **func3**.
- **imm** - immediate numbers. These value can be scrambled within the instruction encoding. For example, in SB, the 11st bit of the actual value was encoded at the 7th bit of the instruction, while the 12rd bit was encoded at the 31th bit.

An RISK-XVII program can be illustrated as below:

```
[Instruction 1 (32 bits)]
[Instruction 2 (32 bits)]
[Instruction 3 (32 bits)]

[...]

[Instruction n (32 bits)]
```

RISK-XVII Instructions

We will now cover in detail all instructions defined in RISK-XVII. Pay close attention as your virtual machine need to be able to decode and execute all of these to be eligible for a full mark! **You are encouraged to summarise them into a reference table before implementing your code.**

Let's use M to denote the memory space. $M[i]$ denotes access to the memory space using address i . For example, to write an immediate value 2017 to the first word of data memory: $M[0 \times 400] = 2017$. Similarly, we use R to denote the register bank, e.g. $R[rd] = R[rs1] + R[rs2]$ denotes an operation that adds the value in $rs1$ and $rs2$, then store the result into rd .

Arithmetic and Logic Operations

1 **add**

- Add - This instruction simply adds two numbers together.
- Operation - $R[rd] = R[rs1] + R[rs2]$
- Encoding:
 - Type: R
 - opcode = 0110011
 - func3 = 000
 - func7 = 0000000

2 **addi**

- Add Immediate - Add a number from register with an immediate number.
- Operation - $R[rd] = R[rs1] + imm$
- Encoding:
 - Type: I
 - opcode = 0010011
 - func3 = 000

3 **sub**

- Subtract
- Operation - $R[rd] = R[rs1] - R[rs2]$
- Encoding:
 - Type: R
 - opcode = 0110011
 - func3 = 000
 - func7 = 0100000

4 lui

- Load Upper Immediate - Load the upper part of an immediate number into a register and set the lower part to zeros.
- Operation - $R[rd] = \{31:12 = \text{imm} \mid 11:0 = 0\}$
- Encoding:
 - Type: U
 - opcode = 0110111

5 xor

- XOR
- Operation - $R[rd] = R[rs1] \wedge R[rs2]$
- Encoding:
 - Type: R
 - opcode = 0110011
 - func3 = 100
 - func7 = 0000000

6 xori

- XOR Immediate
- Operation - $R[rd] = R[rs1] \wedge \text{imm}$
- Encoding:
 - Type: I
 - opcode = 0010011
 - func3 = 100

7 or

- OR
- Operation - $R[rd] = R[rs1] \mid R[rs2]$
- Encoding:
 - Type: R
 - opcode = 0110011
 - func3 = 110
 - func7 = 0000000

8 ori

- OR Immediate
- Operation - $R[rd] = R[rs1] \mid imm$
- Encoding:
 - Type: I
 - opcode = 0010011
 - func3 = 110

9 and

- AND
- Operation - $R[rd] = R[rs1] \& R[rs2]$
- Encoding:
 - Type: R
 - opcode = 0110011
 - func3 = 111
 - func7 = 0000000

Assignment Project Exam Help

<https://tutorcs.com>

10 andi

- AND Immediate
- Operation - $R[rd] = R[rs1] \& imm$
- Encoding:
 - Type: I
 - opcode = 0010011
 - func3 = 111

WeChat: cstutorcs

11 sll

- Shift Left
- Operation - $R[rd] = R[rs1] \ll R[rs2]$
- Encoding:
 - Type: R
 - opcode = 0110011
 - func3 = 001
 - func7 = 0000000

12 srl

- Shift Right
- Operation - $R[rd] = R[rs1] \gg R[rs2]$
- Encoding:
 - Type: R
 - opcode = 0110011
 - func3 = 101
 - func7 = 0000000

13 sra

- **Rotate Right** - the right most bit is moved to the left most after shifting.
- Operation - $R[rd] = R[rs1] \ggg R[rs2]$
- Encoding:
 - Type: R
 - opcode = 0110011
 - func3 = 101
 - func7 = 0100000

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Memory Access Operations**14 lb**

- Load Byte - Load a 8-bit value from memory into a register, and sign extend the value.
- Operation - $R[rd] = \text{sext}(M[R[rs1] + \text{imm}])$
- Encoding:
 - Type: I
 - opcode = 0000011
 - func3 = 000

15 lh

- Load Half Word - Load a 16-bit value from memory into a register, and sign extend the value.
- Operation - $R[rd] = \text{sext}(M[R[rs1] + \text{imm}])$
- Encoding:
 - Type: I
 - opcode = 0000011
 - func3 = 001

16 lw

- Load Word - Load a 32-bit value from memory into a register
- Operation - $R[rd] = M[R[rs1] + \text{imm}]$
- Encoding:
 - Type: I
 - opcode = 0000011
 - func3 = 010

Assignment Project Exam Help

<https://tutorcs.com>

17 lbu

- Load Byte Unsigned - Load a 8-bit value from memory into a register
- Operation - $R[rd] = M[R[rs1] + \text{imm}]$
- Encoding:
 - Type: I
 - opcode = 0000011
 - func3 = 100

WeChat: cstutorcs

18 lhu

- Load Half Word Unsigned - Load a 16-bit value from memory into a register
- Operation - $R[rd] = M[R[rs1] + \text{imm}]$
- Encoding:
 - Type: I
 - opcode = 0000011
 - func3 = 101

19 sb

- Store Byte - Store a 8-bit value to memory from a register.
- Operation - $M[R[rs1] + imm] = R[rs2]$
- Encoding:
 - Type: S
 - opcode = 0100011
 - func3 = 000

20 sh

- Store Half Word - Store a 16-bit value to memory from a register.
- Operation - $M[R[rs1] + imm] = R[rs2]$
- Encoding:
 - Type: S
 - opcode = 0100011
 - func3 = 001

21 sw

- Store Word - Store a 32-bit value to memory from a register.
- Operation - $M[R[rs1] + imm] = R[rs2]$
- Encoding:
 - Type: S
 - opcode = 0100011
 - func3 = 010

Program Flow Operations**22 slt**

- Set if Smaller - Set rd to 1 if the value in rs1 is smaller than rs2, 0 otherwise.
- Operation - $R[rd] = (R[rs1] < R[rs2]) ? 1 : 0$
- Encoding:
 - Type: R
 - opcode = 0110011
 - func3 = 010
 - func7 = 0000000

23 slti

- Set if Smaller Immediate - Set `rd` to 1 if the value in `rs1` is smaller than `imm`, 0 otherwise.
- Operation - $R[rd] = (R[rs1] < imm) ? 1 : 0$
- Encoding:
 - Type: I
 - opcode = 0010011
 - func3 = 010
 - func7 = 0000000

24 sltu

- Set if Smaller - Set `rd` to 1 if the value in `rs1` is smaller than `rs2`, 0 otherwise. Treat numbers as unsigned.
- Operation - $R[rd] = (R[rs1] < R[rs2]) ? 1 : 0$
- Encoding:
 - Type: R
 - opcode = 0110011
 - func3 = 011
 - func7 = 0000000

25 sltiu

- Set if Smaller Immediate - Set `rd` to 1 if the value in `rs1` is smaller than `imm`, 0 otherwise. Treat numbers as unsigned.
- Operation - $R[rd] = (R[rs1] < imm) ? 1 : 0$
- Encoding:
 - Type: I
 - opcode = 0010011
 - func3 = 011
 - func7 = 0000000

26 beq

- Branch if Equal.
- Operation - `if (R[rs1] == R[rs2]) then PC = PC + (imm « 1)`
- Encoding:
 - Type: SB
 - opcode = 1100011
 - func3 = 000

27 bne

- Branch if Not Equal.
- Operation - `if (R[rs1] != R[rs2]) then PC = PC + (imm « 1)`
- Encoding:
 - Type: SB
 - opcode = 1100011
 - func3 = 001

Assignment Project Exam Help

<https://tutorcs.com>

28 blt

- Branch if Less Than/
- Operation - `if (R[rs1] < R[rs2]) then PC = PC + (imm « 1)`
- Encoding:
 - Type: SB
 - opcode = 1100011
 - func3 = 100

WeChat: cstutorcs

29 bltu

- Branch if Less Than. Treat numbers as unsigned.
- Operation - `if (R[rs1] < R[rs2]) then PC = PC + (imm « 1)`
- Encoding:
 - Type: SB
 - opcode = 1100011
 - func3 = 110

30 bge

- Branch if Greater Than or Equal.
- Operation - $\text{if}(R[rs1] \geq R[rs2]) \text{ then } PC = PC + (\text{imm} \ll 1)$
- Encoding:
 - Type: SB
 - opcode = 1100011
 - func3 = 101

31 bgeu

- Branch if Greater Than or Equal. Treat numbers as unsigned.
- Operation - $\text{if}(R[rs1] \geq R[rs2]) \text{ then } PC = PC + (\text{imm} \ll 1)$
- Encoding:
 - Type: SB
 - opcode = 1100011
 - func3 = 111

32 jal

- Jump and Link - Jump to target code address, and save the PC value of next instruction into a register.
- Operation - $R[rd] = PC + 4; PC = PC + (\text{imm} \ll 1)$
- Encoding:
 - Type: UJ
 - opcode = 1101111

33 jalr

- Jump and Link Register - Jump to target code address from a register, and save the PC value of next instruction into a register.
- Operation - $R[rd] = PC + 4; PC = R[rs1] + \text{imm}$
- Encoding:
 - Type: I
 - opcode = 1100111
 - func3 = 000

Virtual Routines

Virtual routines are operations mapped to specific memory addresses such that a memory access operation at that address will have different effects. This can be used to allow programs running in the virtual machine to communicate with the outside world through input/output (I/O) operations.

As part of your task to implement necessary I/O functions for your virtual machine, you are required to develop the following routines:

1 0x0800 - Console Write Character

A memory store command to this address will cause the virtual machine to print the value being stored as a single ASCII encoded character to **stdout**.

2 0x0804 - Console Write Signed Integer

A memory store command to this address will cause the virtual machine to print the value being stored as a single 32-bit signed integer in decimal format to **stdout**.

3 0x0808 - Console Write Unsigned Integer

A memory store command to this address will cause the virtual machine to print the value being stored as a single 32-bit unsigned integer in lower case hexadecimal format to **stdout**.

4 0x080C - Halt

A memory store command to this address will cause the virtual machine to halt the current running program, then output CPU Halt Requested to **stdout**, and exit, regardless the value to be stored.

5 0x0812 - Console Read Character

A memory load command to this address will cause the virtual machine to scan input from **stdin** and treat the input as an ASCII-encoded character for the memory load result.

6 0x0816 - Console Read Signed Integer

A memory load command to this address will cause the virtual machine to scan input from **stdin** and parse the input as a signed integer for the memory load result.

7 0x0820 - Dump PC

A memory store command to this address will cause the virtual machine to print the value of PC in lower case hexadecimal format to **stdout**.

8 0x0824 - Dump Register Banks

A memory store command to this address will force the virtual machine to perform an Register Dump. See Error Handling.

9 0x0828 - Dump Memory Word

A memory store command to this address will cause the virtual machine to print the value of $M[v]$ in lower case hexadecimal format to **stdout**. v is the value being stored interpreted as an 32-bit unsigned integer.

10 0x0830, 0x0834 - Heap Banks

These addresses are used by a hardware-enabled memory allocation system.

11 0x0850 and above - Reserved

Our program will not call these addresses. You are free to use them for your own test cases and debugging purposes.

Heap Banks

<https://tutorcs.com>

One of the biggest selling points of our RISK-XVII-based virtual machine is the hardware (virtual) enabled memory allocation support in addition to the two built-in static memory banks. Programs running inside the virtual machine can request more memory by interfacing with the allocation sub-system through virtual routines.

Your virtual machine program should manage memory allocation requests and ensure that the ownership of each block is always unique to `malloc` requests unless it is not used (free). You need to manage a total of 128 memory banks, each with 64 bytes. Each memory bank is a memory device that can be accessed as a linear array of bytes. To handle allocation requests larger than the size of a single bank, multiple consecutive banks need to be searched and allocated for the request. An error is returned if it is not possible to fulfill such a memory request. The mapped address of the initial byte of the first bank is 0xb700, and there are a total of 8192 bytes of memory that can be dynamically allocated.

Specification The below interfaces are defined for the virtual machine program:

1 0x0830 - malloc

A memory store command to this address will request a chunk of memory with the size of the value being stored to be allocated. The pointer of the allocated memory (starting address) will be stored in $R[28]$. If the memory cannot be allocated, $R[28]$ should be set to zero.

2 0x0834 - free

A memory store command to this address will free a chunk of memory starting at the value being stored. If the address provided was not allocated, an illegal operation error should be raised.

Example Let's consider a scenario where all 128 banks are not allocated yet (free), and we need to handle a `malloc` request with size 270. To fulfill the request, we need to find five free banks that are located consecutively, e.g., `64 + 64 + 64 + 64 + 14`. We will set the first five blocks as used and return the address at the beginning of the first block: `R[28] = &Block[0]`.

Now suppose another request just came in with size 12. We need only a single block to fulfill the request, and after a search the first one-consecutive block is the sixth. We will mark the sixth block as used and return the address to complete the allocation: `R[28] = Block[5]`.

Hint: You are encouraged to use a linked list internally to store and maintain a record of the current allocation.

Error Handling **Assignment Project Exam Help**

Register Dump

<https://tutorcs.com>

When an register dump was requested, the virtual machine should print the value of all registers, including `PC`, in lower case hexadecimal format to **stdout**. The output should be one register per line in the following format:

WeChat: cstutorcs

```
PC = 0x00000001;
R[0] = 0xffffffff;
R[1] = 0xffffffff;

R[...] = 0xffffffff;

R[31] = 0xffffffff
```

Not Implemented

If an unknown instruction was detected, your virtual machine program should output `Instruction Not Implemented:` and the hexadecimal value of the encoded instruction to `stdout`, followed by a **Register Dump** and terminate. For example, when if `0xffffffff` were found in the instruction memory, your program should output:

```
Instruction Not Implemented: 0xffffffff
```

Illegal Operation

When an illegal operation was raised, your virtual machine program should output `Illegal Operation:` and the hexadecimal value of the encoded instruction to `stdout`, followed by a **Register Dump** and terminate.

Any memory accesses outside of defined boundaries will cause this error to be raised.

Starting the Virtual Machine

A binary file will be provided as an image of the instruction and data memory at the beginning of the program. This file can be found by opening the file path supplied as the first command line argument. The below C code outlines the format of the binary input file as a `struct`:

```
#define INST_MEM_SIZE 1024
#define DATA_MEM_SIZE 1024
struct blob {
    char [INST_MEM_SIZE] inst_mem;
    char [DATA_MEM_SIZE] data_mem;
}
```

All registers, including PC, will be initialised to zero at the initialisation of the virtual machine. During each cycle, the virtual machine should fetch and execution the instruction pointed by PC, and increase PC by 4.

Example 1 - Printing "H"

The following RISK-XVII assembly program will print the first letter of "Hello, World" to `stdout`:

```
00000000 <_start>:
    0:    7ff00113      addi    sp,x0,2047
    4:    004000ef      jal     ra,8 <main>
    8:    000017b7      lui     a5,0x1
    c:    80078623      sb      zero,-2036(a5)

00000008 <main>:
    10:    000017b7      lui     a5,0x1
    14:    04800713      addi    a4,x0,72
    18:    80e78023      sb      a4,-2048(a5)
    1c:    00000513      addi    a0,x0,0
    20:    00008067      ret
```

In the program above, the `<_start>` subroutine will initialize the virtual machine to prepare for the C runtime. `ret` is a mnemonic for the instruction `jalr` where `rs1 = 1`. The `start` function loads the address of C stack to `R[2]`, aka `sp`, the stack pointer using the `addi` instruction. This

is necessary to setup a runtime for the machine to execute C program. Using `addi` with the zero register as the source operand is a common method to load immediate value into a register.

Below is the source code for `<_start>` used in the example:

```
_start:

    # init stack pointer
    lui sp, %hi(__stack_end)
    add sp, sp, %lo(__stack_end)

    # jump to main
    call main

    # call halt routine.
    lui      a5, 0x1
    sb       zero, -2036(a5)
```

The stack address was initialised to the bottom of data memory which indicates that the stack is empty, as stack grows upwards. `sb` instruction writes to `0x0800` which will call the `ConsoleWriteCharacter` virtual routine to print of the character.

The second column indicates encoded instruction in hexadecimal format. Below illustrates the memory image input file for this program.

```
[1301F07F]
[EF004000]
[B7170000]
[23860780]
[B7170000]
[13078004]
[2380E780]
[13050000]
[67800000]
[988 bytes of padding for instruction memory]
[1024 bytes of padding for data memory]
```

The input memory image binary file will always be **2048** bytes large.

Below is the equivalent C code to the RISK-XVII assembly program above:

```
char volatile *const ConsoleWriteChar = (char *)0x0800;

int main() {
    *ConsoleWriteChar = 'H';
    return 0;
}
```

Example 2 - Adding two numbers

The following C program running in the virtual machine will scan for two signed integer in `stdin`, then print the sum of these numbers to `stdout`:

```
int volatile *const ConsoleWriteSInt = (int *)0x0804;

inline int scan_char() {
    int result;
    int addr = 0x0816;

    asm volatile("lw %[res], 0(%[adr])"
                  : [res]"=r"(result)
                  : [adr]"r"(addr));

    return result;
}

int main() {
    int a = scan_char();
    int b = scan_char();

    *ConsoleWriteSInt = a + b;

    return 0;
}
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

And the equivalent RISK-XVII assembly program:

```
00000000 <_start>:
    0:      7ff00113          li      sp,2047
    4:      004000ef          jal     ra,8 <main>
    8:      000017b7          lui     a5,0x1
    c:      80078623          sb      zero,-2036(a5)

00000008 <main>:
    10:     00001737          lui     a4,0x1
    14:     81670793          addi    a5,a4,-2026
    18:     0007a683          lw      a3,0(a5)
    1c:     0007a783          lw      a5,0(a5)
    20:     00d787b3          add     a5,a5,a3
    24:     80f72223          sw      a5,-2044(a4)
    28:     00000513          li      a0,0
    2c:     00008067          ret
```

Compilation and Execution

Your virtual machine program will be compiled by running the default rule of a make file. Upon compiling your program should produce a single `vm_riskxvii` binary. Your binary should accept a single argument in the form of the path to a 'RISK-XVII' memory image binary file to execute.

```
make
./vm_riskxvii <memory_image_binary>
```

Please make sure the above commands will compile and run your program. An example Makefile has been provided in the Scaffold, but you're encouraged to customize it to your needs. Additionally, consider implementing the project using multiple C source files and utilizing header files.

Tests will be compiled and run using two make rules; `make tests` and `make run_tests`.

```
make tests
make run_tests
```

These rules should build any tests you need, then execute each test and report back on your correctness.

Failing to adhere to these conventions will prevent your markers from running your code and tests. In this circumstance you will be awarded a mark of 0 for this assignment.

Marking Criteria

The following is the marking breakdown, each point contributes a portion to the total 5% + 30% of the assignment. You will receive a result of zero if your program fails to compile.

For full marks the total size on disk of your program's binary should not exceed 20kB.

Marks are allocated on the basis of:

- Automated Test Cases - 5 - Passing automatic test cases, a number of tests will *not* be released or run until after your final submission.
- Viva - 30 - You will need to answer questions from a COMP2017 teaching staff member regarding your implementation. You will be required to attend a zoom session with COMP2017 teaching staff member after the code submission deadline. A reasonable attempt will need to be made, otherwise you will receive zero for the assessment.

In this session, you will be asked to explain:

- How your program counter is affected by the opcodes executed.
- How your code organises and manages the stack memory for function calls.
- What are the edge cases you considered for when your program returns 1.
- Answer further questions.

- Your code will also be assessed on C coding style conventions (see Ed resources). Clean code will attract the best grade.

Correctness is based on:

- `vm_riskxvii` - return value of the program.
- `vm_riskxvii` - printed standard output of the binary program execution (all virtual routines).

Additionally marks will be deducted on the basis of:

- **Compilation** - If your submission does not compile you will receive an automatic mark of zero for this assessment.
- **Style** - Poor code readability will result in the deduction of marks. Your code and test cases should be neatly divided between header and source files in appropriate directories, should be commented, contain meaningful variable names, useful indentation, white space and functions should be used appropriately. Please refer to this course's style guide for more details.
- **Tests** - A lack of tests, or a lack of thorough testing will result in the deduction of marks. Please provide your test cases along with appropriate scripts to build, run and report the results of your tests. As a number of tests will not be released until after your final submission you are strongly encouraged to test all aspects of your program.
- **Graceful Error Handling** - The description above contains a number of undefined behaviours. Your program should gracefully catch each of these error, report, and exit. Should your program crash marks will be deducted. Your error messages should be meaningful in addition to the specified error handling format.
- **Memory Leaks** - Code that leaks memory will receive a mark of 0.

Warning: Any attempts to deceive or disrupt the marking system will result in an immediate zero for the entire assignment. Negative marks can be assigned if you do not follow the assignment description or if your code is unnecessarily or deliberately obfuscated.

Academic declaration

By submitting this assignment you declare the following:

I declare that I have read and understood the University of Sydney Student Plagiarism: Coursework Policy and Procedure, and except where specifically acknowledged, the work contained in this assignment/project is my own work, and has not been copied from other sources or been previously submitted for award or assessment.

I understand that failure to comply with the Student Plagiarism: Coursework Policy and Procedure can lead to severe penalties as outlined under Chapter 8 of the University of Sydney By-Law 1999 (as amended). These penalties may be imposed in cases where any significant portion of my submitted work has been copied without proper acknowledgment from other sources, including published works, the Internet, existing programs, the work of other students, or work previously submitted for other awards or assessments.

I realise that I may be asked to identify those portions of the work contributed by me and required to demonstrate my knowledge of the relevant material by answering oral questions or by undertaking supplementary work, either written or in the laboratory, in order to arrive at the final assessment mark.

I acknowledge that the School of Computer Science, in assessing this assignment, may reproduce it entirely, may provide a copy to another member of faculty, and/or communicate a copy of this assignment to a plagiarism checking service or in-house computer program, and that a copy of the assignment may be maintained by the service or the School of Computer Science for the purpose of future plagiarism checking.